

GEOMETRY BY CONSTRUCTION: A GEOMETRIC PRIOR BOTTLENECK

Anonymous authors

Paper under double-blind review

ABSTRACT

The conditional entropy bottleneck (CEB) compresses representations against a label-conditional reference whose geometry is left entirely to the optimizer. We show that this reference has two failure channels: in deterministic settings the CEB knob is dead (a known Caveat), and, more fundamentally, the certificate is blind to the relationships between labels: no cross-class term constrains the geometry among the class priors. We close this channel by declaring the geometry: the prior is confined to a qualified geometric family computed statelessly at every training step: an orthogonal class frame (OPB) for classification and an isometric regression axis (EPB) for regression, both with prediction rules tied to the geometry by construction (the zero-parameter anchor-energy classifier and the tied projection head). The framework preserves the CEB certificate, transfers the prior’s separation to the posterior under β -controlled anchor-alignment bounds, and removes the unspecified inter-class geometry of the reference. Across nine datasets and four backbones, GPB is best or tied-best in 7 of 12 settings and dominates CEB everywhere, and the anchor scale converts the β -knob from a collapse dial into a plateau dial.

1 INTRODUCTION

The information bottleneck (IB) principle (Tishby et al., 2000) seeks a representation Z of an input X that retains the information needed to predict the label Y while discarding the rest. The variational IB (VIB) (Alemi et al., 2017) bounds the compression term with a fixed marginal prior $\mathcal{N}(0, I)$; the conditional entropy bottleneck (CEB) (Fischer, 2020) sharpens the bound to the residual information by replacing the marginal prior with a trainable label-conditional reference $b(z|y)$, certified by the gap $\text{ReX} = \mathbb{E}[\text{KL}(q(z|x) \| b(z|y))] \geq I(X; Z|Y)$. We call this gap the *certificate*; its failure modes are our starting point.

The problem: a reference with no declared geometry. Two facts about CEB are established. First, *the knob is dead*: in the deterministic setting $Y = f(X)$, every compression weight selects the same corner of the information plane (Kolchinsky & Tracey, 2019); this is settled background, repaired by their squared compression term, and we return to it only as a mechanism connection (Proposition A.4, Appendix). Second, and less recognized, *the certificate is blind to the relationships between labels*: the CEB compression term is a sum of per-class matching terms (each sample’s KL touches only its own class prior, and no term couples two distinct class priors), so at its class-wise matched optimum it reduces to $I(X; Z|Y)$, irrespective of the inter-class geometry realized by those matched conditionals (Proposition 2.1). Class means may merge, coincide, or point anywhere without changing the certificate; on the family $T_\alpha = B_\alpha \cdot Y$ that interpolates from collapse to the exact copy of the label, the certificate reads zero across the whole retention axis (Proposition A.2, Appendix). CEB therefore compresses against a reference whose inter-class geometry the optimizer may choose at will: what a properly compressed representation should look like (class means separated, ordered, or laid out on a frame) is exactly what the certificate leaves unspecified.

A geometric prior bottleneck. We close this gap by *declaring* the geometry of the reference. The prior $p(z|y)$ is confined to a *qualified geometric family*: label-conditional Gaussians with a bounded diagonal covariance and a mean map that encodes the label structure at a fixed, declared scale: a two-sided separation condition for discrete labels and a two-sided metric condition for continuous labels (Section 2.1). The prior is recomputed statelessly at every step, so the degenerate references

— merged classes, a collapsed point, an arbitrarily rescaled geometry — are excluded by construction rather than by the optimizer’s good behavior. We call this family of objectives a *geometric prior bottleneck* (GPB). Two instantiations realize the principle: *OPB* (orthogonal prior bottleneck) for classification, where the prior means form the scaled orthogonal frame $a q_{\theta,k}$ (every pair of distinct classes at distance $\sqrt{2}a$, with class-mean Gram matrix $a^2 I_K$ by construction), and *EPB* (equidistant prior bottleneck) for regression, where the prior means lie on a unit axis, $\rho \tilde{y} u$, so latent distances equal ρ times standardized label distances exactly. Both keep the CEB posterior, a single KL, and the residual certificate unchanged; the classification head is the anchor-energy classifier $\propto \exp(-\frac{1}{2\tau^2} \|z - a q_{\theta,k}\|^2)$ and the regression head the tied projection $\hat{y} = u^\top z / \rho$: prediction rules tied to the prior geometry by construction, with no free parameters. The constraints bind the *reference*, not the deployed representation; the framework constrains the geometric *form* of the label code and does not claim that the certificate meters retained label information.

Guarantees. The framework’s theory is one chain. The certificate gap identity holds pointwise for any stateless prior (Proposition 2.3), so CEB’s residual bound transfers. And at any ε -near-optimum of the GPB objective, the expected anchor mismatch is $O(1/\beta)$: $\mathbb{E}[D(Y)] \leq (C_{\text{align}} + \varepsilon)/\beta$, with $C_{\text{align}} \leq \ln K$ in the deterministic K -class setting and $C_{\text{align}} = \tau^2/\rho^2$ for the isometric axis (Theorem 2.1, parts (a) and (b)). Combined with a triangle-inequality transfer, the posterior class centers inherit the prior’s separation: $\|m_i - m_j\| \geq \|m_\theta(i) - m_\theta(j)\|$ minus mismatch terms of order $1/\sqrt{\beta}$ (Proposition 2.4 and its corollaries). The framework pins exactly what CEB left free: the geometry of the reference.

Experiments. Across nine datasets and four backbones, GPB is best or tied-best in 7 of 12 settings and dominates CEB in all 12; the mechanism-level validation confirms the story end to end: the constructed geometry transfers to the posterior (nearly orthogonal class centers on classification, an on-axis posterior on regression), and prediction survives with the zero-parameter geometric heads, matching or exceeding the free classifier.

Our contributions are: (1) the identification and formalization of the missing cross-label geometric constraint in CEB (Proposition 2.1); (2) the GPB framework and its two instances, OPB and EPB, with prediction rules tied to the prior geometry; (3) a unified alignment-to-separation guarantee: certificate preservation and β -controlled conditional separation transfer for any qualified geometric prior, instantiated by the orthogonal frame and the isometric axis; and (4) mechanism-level empirical evidence that the constructed geometry is transferred to the posterior and is sufficient for prediction.

2 METHOD

2.1 FROM CEB TO GPB

Let X be the input and Y the label with K classes; the encoder induces the Markov chain $Y - X - Z$. Following Fischer (2020), the variational CEB objective is $\min_{e,b,c} \langle \log e(z|x) \rangle - \langle \log b(z|y) \rangle - \gamma \langle \log c(y|z) \rangle$ over the stochastic encoder e , the backward encoder b , and the classifier c . We use the diagonal Gaussian posterior $q_\phi(z|x) = \mathcal{N}(\mu_\phi(x), \text{diag}(\sigma_\phi^2(x)))$, the loss convention $\text{CE} + \beta \text{KL}$ ($\beta = 1/\gamma$), and the deterministic code $z = \mu_\phi(x)$ at evaluation. The central object is the residual certificate: for any conditional prior $r(z|y)$,

$$\mathbb{E}[\text{KL}(q_\phi(z|x) \| r(z|y))] = I(X; Z|Y) + \mathbb{E}_y[\text{KL}(q_\phi(z|y) \| r(z|y))] \geq I(X; Z|Y), \quad (1)$$

by Fischer’s chain-rule identity. CEB instantiates $r(z|y) = b_\theta(z|y)$; VIB fixes $r = \mathcal{N}(0, I)$. The certificate’s failure mode is that its compression term carries no cross-class structure:

Proposition 2.1 (The certificate records nothing about inter-label geometry). *The CEB compression term contains no explicit term coupling two distinct class priors: it is a sum of per-class matching terms, each sample’s KL touching only its own class prior. Consequently, at its class-wise matched optimum ($b = q(z|y)$), the certificate reduces to*

$$\text{ReX} = I(X; Z|Y), \quad (2)$$

irrespective of the inter-class geometry realized by the matched conditionals: class means merged, coincident, or placed in arbitrary directions all yield the same certificate value. (The prediction term couples classes through the softmax; the statement is at a fixed encoder with a matched backward encoder.)

At the matched optimum the certificate thus carries no information about the relationships among the class priors, and on the retention family $T_\alpha = B_\alpha \cdot Y$ of Kolchinsky & Tracey (2019) it reads zero from total collapse to the full copy (Proposition A.2, Appendix). We close this channel by *declaring* the geometry of the reference.

Let $\mathcal{G}$ be a *qualified geometric prior family*: a set of label-conditional distributions $p_\theta(z|y)$ that depend on y alone, computed from the prior encoder, with a bounded diagonal covariance,

$$p_\theta(z|y) = \mathcal{N}(m_\theta(y), \Sigma_\theta(y)), \quad \Sigma_\theta(y) \text{ diagonal}, \quad \Sigma_\theta(y) \preceq \tau_{\max}^2 I_d \text{ for all } y, \quad (3)$$

and a mean map that encodes the label structure at a fixed, declared scale: for discrete labels a two-sided separation condition,

$$\Delta \leq \|m_\theta(y) - m_\theta(y')\| \leq \bar{\Delta} \text{ for all } y \neq y', \quad (4)$$

and for continuous labels with a declared label metric d_Y ,

$$\rho_{\min} d_Y(y, y') \leq \|m_\theta(y) - m_\theta(y')\| \leq \rho_{\max} d_Y(y, y'). \quad (5)$$

The declared constants $\Delta, \bar{\Delta}, \rho_{\min}, \rho_{\max}$ are fixed before training and not trained. The lower bounds encode the label distinctions; the upper bounds close the scale escape: without them the family would contain $m_\theta(y) = s m_0(y)$ for arbitrarily large s , and the objective might lack a finite optimum, making “geometry by construction” vacuous. The GPB objective is

$$\mathcal{L}_{\text{GPB}} = \mathbb{E}_{(x,y)} \left[\text{CE}(c_\psi(z), y) + \beta \text{KL}(q_\phi(z|x) \| p_\theta(z|y)) \right], \quad z \sim q_\phi(z|x), \quad p_\theta \in \mathcal{G}. \quad (6)$$

2.2 GEOMETRIC INSTANTIATIONS

2.2.1 OPB: AN ORTHOGONAL CLASS FRAME (CLASSIFICATION)

For classification, let $K \leq d$. The prior encoder g_θ is a single unactivated linear map, evaluated on the all-class one-hot matrix E_K at every forward pass, giving the raw class-mean matrix

$$U_\theta = g_\theta(E_K)^\top \in \mathbb{R}^{d \times K}, \quad (7)$$

orthonormalized by the thin polar factor

$$Q_\theta = U_\theta (U_\theta^\top U_\theta)^{-1/2}, \quad Q_\theta^\top Q_\theta = I_K, \quad (8)$$

computed via the thin SVD $U_\theta = W \Sigma V^\top$, $Q_\theta = W V^\top$ (no sign convention needed). The polar factor is exactly equivariant under relabeling of the classes, $Q_\theta(U_\theta P) = Q_\theta(U_\theta) P$ (derivation in Appendix A.3), so no class index plays a distinguished role. We make one standing assumption, verified, not imposed, in the implementation:

Assumption 2.1 (Full rank). *U_θ has full column rank at every training step, so that Eq. equation 8 and its backward pass are defined. The condition holds almost surely at the identity initialization, and $\sigma_{\min}(U_\theta)$ is monitored along training (Appendix A.3).*

The geometric prior replaces the raw means by the scaled frame

$$p_\theta^{\text{OPB}}(z|y=k) = \mathcal{N}(a q_{\theta,k}, \tau^2 I_d), \quad q_{\theta,k} := Q_\theta[:, k], \quad (9)$$

with anchor scale $a > 0$ and fixed shared prior variance τ^2 (default 1.0): a learnable per-class variance would give the reference an extra channel to absorb the compression cost by inflating its covariance, precisely the degeneracy the construction excludes. Since $Q_\theta^\top Q_\theta = I_K$, the class-mean Gram matrix is $a^2 I_K$ by construction:

$$a_i^\top a_j = a^2 \delta_{ij}, \quad \|a_k\| = a, \quad \|a_i - a_j\|^2 = 2a^2 \quad (i \neq j). \quad (10)$$

The frame is stateless: an immediate function of the current prior-encoder parameters, computed from the all-class table: no EMA, no prototypes, no batch statistics. The single-KL objective is

$$\mathcal{L}_{\text{OPB}} = \mathbb{E}_{(x,y)} \left[\text{CE}(c_\psi(z), y) + \beta \text{KL}(q_\phi(z|x) \| p_\theta^{\text{OPB}}(z|y)) \right], \quad z \sim q_\phi(z|x). \quad (11)$$

The main classifier is the anchor-energy classifier, whose scale is fixed by construction:

Proposition 2.2 (Anchor-energy classifier equivalence). *The scores*

$$s_k(z) = -\frac{\|z - a q_{\theta,k}\|^2}{2\tau^2} + \log \pi_k \quad (12)$$

are the class- k Gaussian log posterior up to a class-independent term; for uniform classes the equivalent linear classifier has weights $w_k = (a/\tau^2) q_{\theta,k}$: equal norms and pairwise orthogonal rows. The prediction rule therefore uses exactly the same geometry as the prior, with no free classifier scale (proof in Appendix A.3).

The posterior logvar head is zero-initialized and the prior encoder identity-initialized ($U_\theta = [e_1, \dots, e_K]$, full column rank), so training starts at $\text{KL} \approx a^2/(2\tau^2)$ with a stable polar backward. The plain linear classifier is kept in the experiments as a free-head variant; it does not enjoy the fixed-scale condition of Proposition A.4.

2.2.2 EPB: AN EQUIDISTANT REGRESSION AXIS (REGRESSION)

For a one-dimensional continuous label, standardize once on the training set, $\tilde{y} = (y - y_{\text{center}})/y_{\text{scale}}$, with the statistics frozen for validation and test. Let $W \in \mathbb{R}^{d \times 1}$ be a trainable direction and $u = W/\|W\|$ its per-step normalization. The isometric prior is

$$p^{\text{EPB}}(z|y) = \mathcal{N}(\rho \tilde{y} u, \tau^2 I_d), \quad \|\mu_p(y_i) - \mu_p(y_j)\| = \rho |\tilde{y}_i - \tilde{y}_j|, \quad (13)$$

with isometric scale $\rho > 0$ and fixed shared prior variance τ^2 : latent distances equal ρ times standardized label distances exactly, at every scale. Order and magnitude of label differences are encoded with equality, which qualifies the family (Eq. equation 3 with $\tau_{\text{max}}^2 = \tau^2$ and Eq. equation 5 with $\rho_{\text{min}} = \rho_{\text{max}} = \rho$ on the standardized label metric). The main variant uses the trainable normalized axis and no bias (the fixed-axis and unconstrained-scale ablations are in Appendix A.3). The prediction head is the *tied projection* $\hat{y} = u^\top z/\rho$, the regression analogue of the anchor-energy classifier: geometrically consistent, with no free direction or scale that could bypass the prior; the ordinary regression head is kept as a free-head variant. The objective is $L_{\text{reg}}(\hat{y}, y) + \beta \text{KL}(q_\phi(z|x) \| p^{\text{EPB}}(z|y))$: a single KL, stateless.

2.3 THEORETICAL GUARANTEES

The framework’s theory is one chain: at any ε -near-optimum the expected anchor mismatch is $O(1/\beta)$, and a triangle-inequality transfer then makes the posterior inherit the prior’s geometry. The first link holds pointwise:

Proposition 2.3 (A stateless geometric prior preserves the CEB certificate). *For any prior family $\mathcal{G}$ whose members are conditional priors independent of x , and for every parameter value θ at which $p_\theta \in \mathcal{G}$ is defined, the gap identity of Eq. equation 1 applies pointwise in θ :*

$$\mathbb{E}[\text{KL}(q_\phi(z|x) \| p_\theta(z|y))] = I_{q_\phi}(X; Z|Y) + \mathbb{E}_Y[\text{KL}(q_\phi(z|Y) \| p_\theta(z|Y))] \geq I_{q_\phi}(X; Z|Y). \quad (14)$$

A batch-dependent prior would break the conditional-prior argument.

Proposition 2.4 (Separation transfer for any qualified geometric prior). *Let $\mathcal{G}$ be qualified with variance bound τ_{max}^2 . (a) Discrete labels. With posterior class centers $m_k := \mathbb{E}[\mu_\phi(X)|Y = k]$ and class-wise anchor mismatch*

$$D_k := \frac{1}{2\tau_{\text{max}}^2} \mathbb{E}[\|\mu_\phi(X) - m_\theta(k)\|^2 \mid Y = k], \quad (15)$$

for every pair $i \neq j$,

$$\|m_i - m_j\| \geq \|m_\theta(i) - m_\theta(j)\| - \sqrt{2\tau_{\text{max}}^2 D_i} - \sqrt{2\tau_{\text{max}}^2 D_j} \geq \Delta - \sqrt{2\tau_{\text{max}}^2 D_i} - \sqrt{2\tau_{\text{max}}^2 D_j}, \quad (16)$$

using Eq. equation 4. (b) Continuous labels. For Borel-measurable X, Y with integrable $\mu_\phi(X)$, the regular conditional posterior center $m(y) := \mathbb{E}[\mu_\phi(X) \mid Y = y]$ and pointwise mismatch

$$D(y) := \frac{1}{2\tau_{\text{max}}^2} \mathbb{E}[\|\mu_\phi(X) - m_\theta(y)\|^2 \mid Y = y] \quad (17)$$

are defined for P_Y -a.e. y , and for P_Y -a.e. pair $y \neq y'$,

$$\|m(y) - m(y')\| \geq \|m_\theta(y) - m_\theta(y')\| - \sqrt{2\tau_{\max}^2 D(y)} - \sqrt{2\tau_{\max}^2 D(y')} \geq \rho_{\min} d_Y(y, y') - \sqrt{2\tau_{\max}^2 D(y)} - \sqrt{2\tau_{\max}^2 D(y')}. \quad (18)$$

The second link shows that the objective itself forces the mismatch small, in one theorem stated for the two label-space structures:

Theorem 2.1 (Anchor alignment at near-optima). (a) Deterministic K -class classification. Assume $Y = f(X)$, a cross-entropy prediction loss, a qualified family whose diagonal covariances the posterior can represent, and an encoder family containing class-constant maps, so that the aligned solution $q(z|x) = p_\theta(z|y)$ is feasible; assume further that the classifier family contains the Bayes posterior of the aligned mixture,

$$c_k(z) = \frac{\pi_k |\Sigma_\theta(k)|^{-1/2} e^{-\frac{1}{2}(z - m_\theta(k))^\top \Sigma_\theta(k)^{-1} (z - m_\theta(k))}}{\sum_j \pi_j |\Sigma_\theta(j)|^{-1/2} e^{-\frac{1}{2}(z - m_\theta(j))^\top \Sigma_\theta(j)^{-1} (z - m_\theta(j))}}. \quad (19)$$

Then for any solution (ϕ, θ, ψ) with objective value $L \leq L^* + \varepsilon$, where L^* is the infimum of Eq. equation 6, the expected anchor mismatch is controlled by β :

$$\mathbb{E}[D(Y)] \leq \frac{C_G + \varepsilon}{\beta} \leq \frac{\ln K + \varepsilon}{\beta}, \quad (20)$$

where $C_G := H(Y | Z_{\text{align}})$ is the cross-entropy of the Bayes classifier of Eq. equation 19 on the aligned mixture. In the shared-covariance orthogonal instance the determinant factors cancel and this Bayes posterior reduces exactly to the anchor-energy classifier of Eq. equation 12.

(b) Deterministic continuous regression. Assume $Y = f(X)$ with a non-negative regression loss, a qualified family, a posterior able to realize $q(z|x) = p_\theta(z|y)$, and the tied projection head of the isometric instance. For the standardized squared loss, $\hat{y} = u^\top z / \rho$ with $z \sim p_\theta(z|y)$ gives $\hat{y} = \tilde{y} + (\tau/\rho) \eta$ with $\eta \sim \mathcal{N}(0, 1)$, so the aligned risk is $C_{\text{align}} = \tau^2 / \rho^2$. Then for any solution with $L \leq L^* + \varepsilon$,

$$\mathbb{E}[D(Y)] \leq \frac{C_{\text{align}} + \varepsilon}{\beta}, \quad (21)$$

with $D(y)$ as in Eq. equation 17. (The scope of the deterministic comparison-solution argument is discussed in Appendix A.1.)

Corollary 2.1 (Separation at near-optimal solutions). (a) Classification. Under Theorem 2.1(a), for classes $i \neq j$ with $\pi_i, \pi_j > 0$,

$$\|m_i - m_j\| \geq \|m_\theta(i) - m_\theta(j)\| - \sqrt{\frac{2(C_G + \varepsilon)}{\beta \pi_i}} - \sqrt{\frac{2(C_G + \varepsilon)}{\beta \pi_j}}. \quad (22)$$

(b) Regression. Under Theorem 2.1(b), for any $t > 0$, on the label set $S_t := \{y : D(y) \leq t\}$ of P_Y -measure at least $1 - (C_{\text{align}} + \varepsilon)/(\beta t)$, the pointwise transfer of Proposition 2.4(b) yields, for all $y, y' \in S_t$,

$$\|m(y) - m(y')\| \geq \rho_{\min} d_Y(y, y') - 2\sqrt{2\tau_{\max}^2 t}; \quad (23)$$

for EPB ($\rho_{\min} = \rho$, $\tau_{\max} = \tau$) the bound reads $\rho d_Y(y, y') - 2\sqrt{2\tau^2 t}$.

The orthogonal instance upgrades these statements with a constructive fact:

Proposition 2.5 (Positive rate cost of class-independent means). If the posterior mean is class-independent, $\mathbb{E}[\mu_\phi(X)|Y = k] = c$ for every k , the mean-mismatch part of the rate satisfies

$$\mathbb{E}\left[\frac{1}{2\tau^2} \|\mu_\phi(X) - a q_{\theta, Y}\|^2\right] \geq \frac{a^2}{2\tau^2} \left(1 - \sum_{k=1}^K \pi_k^2\right), \quad (24)$$

which for uniform classes reads $a^2(1 - 1/K)/(2\tau^2)$: the VIB corner $\mu \equiv 0$ is a penalized configuration rather than a free optimum, whenever $a > 0$ and two classes have positive probability.

Corollary 2.2 (The orthogonal instance satisfies the framework separation transfer). *The orthogonal instance is qualified with $\tau_{\max}^2 = \tau^2$ and $\Delta = \bar{\Delta} = \sqrt{2} a$; Proposition 2.4 specializes to*

$$\|m_i - m_j\| \geq \sqrt{2} a - \sqrt{2\tau^2 D_i} - \sqrt{2\tau^2 D_j}, \quad D_k := \frac{1}{2\tau^2} \mathbb{E} \left[\|\mu_\phi(X) - a q_{\theta,k}\|^2 \mid Y = k \right]. \quad (25)$$

Scope. The geometric constraints bind the reference, not the deployed representation: at any frame rotation there exists a full-retention label code with a zero certificate (Remark A.3, Appendix), so the framework constrains the geometric *form* of the label code (precisely what CEB left unspecified), and does not claim that the certificate meters retained label information. The proofs, the dead-knob connection (the anchored reference supplies the quadratic displacement behind the strict convexity of Proposition A.4), and the implementation details are in the appendix.

3 EXPERIMENTS

The experiments address two questions. First, does replacing CEB’s unconstrained reference with a constructed geometry (an orthogonal class frame or an isometric axis) cost or gain anything in end-to-end prediction across heterogeneous settings? Second, does prediction survive when the prediction rule is tied to that geometry by construction (the anchor-energy classifier and the tied projection head), so that the free head’s degrees of freedom are removed? The main sweep answers both on twelve dataset–backbone settings; the compression–accuracy and realized- $\mathbb{E}[\text{KL}]$ curves are reported in Section 3.3, adversarial robustness in Section 3.4, the geometric mechanism validation in Section 3.5, and the remaining mechanism-level checks (nearest-anchor sufficiency, memorization) in the supplementary analysis.

3.1 EXPERIMENTAL SETUP

Datasets. Twelve settings over nine datasets and four backbones, spanning image, graph, and text modalities. *Classification:* MNIST (MLP and CNN backbones); ImageNet-100 (MLP and CNN) on frozen ResNet-50 features (layer-3 activations pooled to 4×4); Cora, transductive node classification (GCN); IMDB sentiment (RNN); AG News (RNN) on frozen BERT layer-10 token features reduced by PCA to 50 dimensions. *Regression:* California Housing (MLP); STS-B (RNN with frozen GloVe-100d embeddings); ZINC-12k molecular graphs (GCN, the official 10k/1k/1k split); AgeDB face age (MLP and CNN). Where no official split exists we use a fixed 60/20/20 train/validation/test partition (seed 42); on MNIST 10k of the 60k training images are held out for validation. Regression targets are standardized to $[0, 1]$ on the training set (the test metric is re-normalized).

Backbones and protocol. The MLP uses hidden layers 512–256; the CNN two convolutional blocks (channels 32, 64) followed by a 256-dim hidden layer; the GCN matches the MLP sizes; the RNN is a two-layer LSTM with 512 and 256 units (a single 256-unit layer with max time-pooling on STS-B). All methods share the backbone, dropout 0.2, and bottleneck dimension $d = 256$. Every run trains with Adam (learning rate 10^{-3} ; 10^{-2} on Cora), at most 100 epochs with early stopping on the validation metric (patience 10), batch size 512 (128 on ZINC), and evaluates the checkpoint with the best validation score on the test set. GPB and GPB-L train five seeds per configuration (seeds 0–4); each variational baseline trains fifteen seeds per β value (seeds 0–14), so that every method trains a matched total budget ($6 \times 3 \times 5 = 6 \times 15 = 90$ runs per method). The tables report the mean over the runs. Classification early-stops on validation macro-AUC and reports test accuracy; regression early-stops on validation R^2 and reports test R^2 .

Baselines. The plain backbone (Base); VIB (Alemi et al., 2017) with the fixed $\mathcal{N}(0, I)$ prior; SVIB, the squared compression term of Kolchinsky & Tracey (2019) ($\text{CE} + \beta \text{KL}^2$); NIB (Kolchinsky et al., 2019), the noise-injection bottleneck with the pairwise-distance upper bound on $I(X; M)$; CEB (Fischer, 2020) with the trainable label-conditional backward encoder; and the supervised β -DVCCA of Abdelaleem et al. (2025), which augments the VIB loss with an input reconstruction term, $\text{CE} + \beta (\text{KL} + \text{MSE}_{\text{recon}})$. All variational methods share the encoder, the posterior heads, and the loss convention of Eq. equation 6.

Table 1: Test accuracy (%) on classification tasks and test R^2 on regression tasks (mean over runs). Each entry reports the best configuration over the tuned β (and anchor-scale) grid; the best entry per dataset is bolded.

Dataset	Backbone	Base	VIB	SVIB	NIB	CEB	DVCCA	GPB	GPB-L
MNIST	MLP	98.33	98.28	98.15	98.27	<u>98.43</u>	98.24	98.54	<u>98.43</u>
MNIST	CNN	99.13	99.28	99.23	99.22	99.30	99.29	99.32	<u>99.31</u>
ImageNet-100	MLP	83.12	82.84	83.02	82.35	82.77	82.99	84.40	<u>84.01</u>
ImageNet-100	CNN	83.40	82.94	83.11	82.91	82.97	83.21	86.08	<u>84.88</u>
Cora	GCN	87.68	86.83	86.46	86.35	86.27	86.90	88.12	<u>87.86</u>
IMDb	RNN	84.85	85.50	85.99	86.14	85.55	85.69	<u>86.84</u>	86.99
AG News	RNN	92.94	92.72	92.86	92.91	92.83	92.78	92.87	<u>92.92</u>
Cal. Housing	MLP	0.7618	<u>0.7901</u>	0.7897	0.7862	0.7894	0.7878	0.7895	0.7914
STS-B	RNN	0.2899	0.2719	0.2697	0.2417	0.2668	0.2713	<u>0.2762</u>	0.2755
ZINC	GCN	0.6348	0.6339	0.6325	0.6178	0.6257	0.6294	<u>0.6344</u>	0.6323
AgeDB	MLP	0.2180	0.3173	<u>0.3385</u>	0.2710	0.3317	0.3180	0.3536	0.3320
AgeDB	CNN	0.3914	0.3342	0.3323	0.3337	0.3882	0.3205	0.4700	0.4388
Mean rank (of 12)		4.25	5.25	5.00	6.50	5.46	5.50	1.75	2.29
Best or tied-best (of 12)		3	0	0	0	0	0	7	2

Our variants. *GPB* denotes our method in the tables. On classification rows it is *OPB*, the orthogonal-frame prior of Eq. equation 9 with the anchor-energy classifier of Eq. equation 12: the prediction rule has no free parameters, its means and variances are exactly the prior’s, and $K \leq d$ holds in all settings (max $K = 100$ on ImageNet-100); on regression rows it is *EPB*, the isometric prior of Eq. equation 13 with the tied projection head $\hat{y} = u^\top z / \rho$. *GPB-L* is the free-head ablation: the identical prior and KL, with an ordinary trainable linear classifier (or regression head) on z , so prediction is free to bypass the constructed geometry. The frame is orthonormalized from the all-class mean matrix at every forward pass by the thin SVD polar factor of Eq. equation 8 ($Q_\theta = WV^\top$); the prior encoder is identity-initialized (initial anchors $e_1, \dots, e_K$, initial KL $\approx a^2/2$ under unit variances), the variance heads are zero-initialized, and the construction is stateless: no EMA, no prototypes, no batch statistics.

Hyperparameter selection. Each method selects its hyperparameters on the validation set from a shared grid: a six-value log-spaced β grid $\{10^{-4}, \dots, 10\}$ for every variational baseline, and, for GPB and GPB-L, an additional three-value anchor-scale grid $\{1, 6, 12\}$ (6×3 combinations). The main table reports, per method and setting, the test metric of the best validation configuration, so all methods are compared under an identical selection protocol (Base has no sweep and early-stops as usual).

3.2 MAIN RESULTS

Table 1 reports test accuracy (classification) and test R^2 (regression), averaged over runs, for each method’s best configuration; bold marks the best mean per row and underline the second best, and the vertical rule separates GPB and its free-head ablation GPB-L from the baselines. Statistical claims (best vs. tied) follow the paired-difference confidence intervals of Appendix A.6. Standard deviations are omitted for width and are reported in Appendix A.5. Three observations.

The constructed geometry is free, and it usually pays. GPB is best or tied-best in 7 of 12 settings with mean rank 1.75, against 4.25 for the plain backbone and 5.0–6.5 for the variational baselines, none of which is ever best. By the paired bootstrap 95% CIs (Appendix A.6), the GPB–CEB difference is significantly positive in 8 of 12 settings and within the non-inferiority margin in three more; only on STS-B does the interval overlap zero beyond the margin ($[-0.009, +0.025]$), so no setting shows a significant CEB advantage. GPB dominates CEB by mean in all 12 settings and is best among all methods on seven rows. The gains concentrate where the label structure is rich or supervision is scarce: $+3.1$ accuracy over CEB on ImageNet-100 with the CNN backbone (86.1 vs. 83.0), $+1.9$ on Cora, and $+0.08$ R^2 on AgeDB over the plain CNN (0.470 vs. 0.388). On saturated settings (MNIST) the margin is small but consistent, and on the three rows where the plain backbone remains best (STS-B, ZINC, AG News) the gap is marginal (≤ 0.03): the geometry is neutral there rather

than harmful, and two of the three are regression tasks where every bottleneck method trails the plain backbone.

Prediction does not require the free head. GPB-L, the free linear head on the same geometric prior, is best or tied-best in 2 of 12 settings (mean rank 2.29), and the GPB-GPB-L gap is within 0.4 accuracy or R^2 points in 11 of 12 settings (the exception is ImageNet-100 CNN, 1.2 points): the zero-parameter anchor-energy classifier — whose weights are the frame itself — reaches the accuracy of a freely trained head in almost all settings, so the degrees of freedom removed by construction are not required for prediction, as claimed in the introduction. The three settings where GPB-L edges out GPB by mean (IMDb, AG News, Cal. Housing) do so with paired CIs overlapping zero, confirming that the tie does not inflate the numbers. On the five regression rows the tied projection head of EPB wins or matches the free regression head in four (the exception is Cal. Housing by 0.002).

The tie regularizes the high-cardinality and low-data regimes. The two settings with the largest GPB margins (ImageNet-100, with $K = 100$, and AgeDB, a small, high-variance regression task) are exactly those where an unconstrained classifier can overfit its own scale or the scarce supervision: fixing the prediction rule to the prior’s geometry removes both escape routes, consistent with the fixed-scale condition under which the aligned surrogate of Proposition A.4 has a unique interior optimum per β .

3.3 COMPRESSION-ACCURACY AND THE KL CERTIFICATE OVER β

The main table selects the best (β, a) per method; the figures below report the compression-accuracy trade-off on one classification and one regression setting (ImageNet-100 and Cal. Housing, both with the MLP backbone). Each setting has two companion panels over the same β grid: the task metric against β (panels a–b), and the realized test-set $\mathbb{E}[\text{KL}]$ — the certificate itself — against β (panels c–d). The certificate panels are what make a β -matched sweep interpretable: the two methods reach different compressions at the same β — on Housing at $\beta = 0.1$, CEB sits at $\mathbb{E}[\text{KL}] \approx 0.09$ while EPB ($\rho = 12$) is at ≈ 0.27 — so the compression each point actually achieves is read off directly rather than assumed from β . No points are omitted: collapsed references and variance explosions are visible at the extremes of the sweep.

Figure 1 exhibits the two roles of the anchor scale across the β sweep, with panels (c)–(d) translating each point into its realized compression. (i) *Without a scale, the reference collapses.* On ImageNet-100, CEB’s accuracy erodes to 71.5% at $\beta = 0.1$ ($\mathbb{E}[\text{KL}] \approx 3.6$) and its reference escapes entirely at $\beta \geq 0.5$ ($\mathbb{E}[\text{KL}] \rightarrow 0$ with accuracy at chance, panel c); OPB at $a = 1$ survives one grid step longer (83.3% at $\beta = 0.1$) and collapses at the same end of the sweep. OPB ($a = 6$) holds accuracy above 74% out to $\beta = 5$ and falls to 64.7% at $\beta = 25$. On Housing, CEB drops to $R^2 = 0.49$ at $\beta = 0.1$ and below 0.1 by $\beta \geq 0.5$, while every EPB curve keeps $R^2 \geq 0.76$ out to $\beta = 25$. Orthogonality alone therefore does not prevent the collapse: the frame is fixed up to rotation, but the anchor-energy classifier’s margin is a^2/τ^2 (Proposition 2.2), and at $a = 1$ that margin is too small to carry a solution at strong compression. (ii) *The anchor scale controls where the collapse sets in, and at matched compression the geometry pays or ties.* On ImageNet-100, OPB ($a = 12$) attains 84.2% at $\beta = 1$ ($\mathbb{E}[\text{KL}] \approx 21$), above CEB’s best point anywhere on its sweep (82.8% at $\beta = 10^{-2}$), and still 79.4% at $\beta = 25$ ($\mathbb{E}[\text{KL}] \approx 2.2$); on Housing the four curves peak within half a point (CEB 0.789, EPB 0.786–0.789), but only EPB retains its peak into the strong-compression regime: the plateau is the flat EPB segment where R^2 stays above 0.76 from $\beta = 0.1$ out to $\beta = 25$ while the certificate keeps shrinking.

A visible cost remains in the weak-compression regime: on ImageNet-100 at $\beta = 10^{-4}$ ($\mathbb{E}[\text{KL}] \gtrsim 300$), OPB trails CEB by one to three points ($a = 6$: 81.0 vs. 82.6; $a = 12$: 79.8 vs. 82.6): the price of the fixed-scale energy rule when compression is weak. On Housing the same regime shows EPB ($\rho = 1$) dipping to $R^2 = 0.53$ at $\beta = 5 \cdot 10^{-5}$, where its certificate explodes to $\approx 10^8$ (panel d): with a unit axis and negligible β the off-axis variance is unconstrained. The construction pays for itself in the strong-compression regime, where CEB is at chance.

The mechanism is the one the theory predicts. At strong β the posterior is pinned near the aligned solution: Theorem 2.1(a) bounds the expected anchor mismatch by $(\ln K + \epsilon)/\beta$, and the accuracy achievable at alignment is the aligned Gaussian mixture’s accuracy, governed by the scale-to-noise ratio a/τ (the nearest-anchor error probability is $\Phi(-a/(\sqrt{2}\tau))$; appendix). With the classifier scale




Figure 1: Compression-accuracy and the realized certificate over β (log scale), $\beta \in \{5 \cdot 10^{-5}, \dots, 25\}$; CEB one curve per panel, GPB three (a on classification, ρ on regression, both $\in \{1, 6, 12\}$). (a)–(b) Test accuracy (%) on ImageNet-100 (MLP) and test R^2 on Cal. Housing (MLP) against β : OPB with $a = 12$ holds the plateau out to $\beta = 25$ while CEB and OPB ($a = 1$) collapse. (c)–(d) The realized test-set $\mathbb{E}[\text{KL}]$ against β (log scale) for the same runs: CEB’s certificate collapses with its reference, EPB ($\rho = 1$) explodes near $\beta = 5 \cdot 10^{-5}$, and the larger anchors keep $\mathbb{E}[\text{KL}]$ in a stable band.

fixed at a/τ^2 by construction, the only remaining lever at full compression is a itself: the anchor scale converts the compression knob from a collapse dial into a plateau dial. The KL decomposition shows the two channels at work: OPB’s $\mathbb{E}[\text{KL}]$ is dominated by the mean-mismatch term (anchor displacement; $a = 12$, $\beta = 1$: 20.6 mean vs. 0.36 variance), while CEB’s reference absorbs compression by shrinking both terms toward zero: the reference chases the posterior, exactly the degeneracy the constructed geometry excludes (Lemma A.1). This is the empirical counterpart of the separation-transfer guarantees (Proposition 2.4; corollaries in the appendix) and of the fixed-scale condition of Proposition A.4; consistent with that proposition’s scope, we do not claim that the map $\beta \mapsto$ accuracy is monotone, only that the constructed scale keeps the plateau from collapsing.

3.4 ADVERSARIAL ROBUSTNESS

We evaluate robustness to white-box attacks following the CEB paper’s *targeted* variant (its Figure 4, bottom panels): targeted projected gradient descent (PGD, 20 steps, step size $2.5\varepsilon/\text{steps}$, random start) toward a single class, with success measured as the fraction of originally correct predictions flipped to that class. The target is digit 0, the most distinctive MNIST class (the analogue of Fashion-MNIST “trouser”), so that class confusion does not masquerade as robustness. ε is defined in the normalized pixel space (MNIST mean/std); we attack 1,000 test images and report the mean and standard deviation over runs. All stochastic models are attacked and evaluated on their *stochastic* path: each PGD step averages gradients over 10 samples of $z \sim q_\phi(z|x)$ (expectation over transformation), and every prediction is the majority class under 10 MC samples. We use $\varepsilon_\infty = 0.3$ and $\varepsilon_2 = 6.0$ — the smallest L_2 budget we tested at which the targeted attack engages the decision boundaries at all (at 2.0 it cannot reach the target region from any other class and sits at a near-zero floor for every model). The sweep covers the full β grid for OPB; CEB is drawn only over its op-




Figure 2: Targeted PGD attack success rate against β (log scale) on MNIST (MLP backbone): target class 0, 20 steps, EOT-10 gradients and 10-sample MC predictions on the stochastic path. (a) L_∞ , $\varepsilon = 0.3$; (b) L_2 , $\varepsilon = 6.0$. The dashed grey line is the deterministic MLP baseline; CEB is drawn over its operating range only ($\beta \leq 0.1$), beyond which its reference collapses to chance. Every curve falls with β , and OPB keeps gaining robustness across the entire grid: at $\beta = 25$, OPB ($a = 1$) reaches 11% — 50 points below the baseline.

erating range $\beta \leq 0.1$, beyond which the reference collapses and accuracy is at chance — a dead model’s zero attack success is not robustness.

Figure 2 shows the targeted-attack picture. Every curve falls as β grows, reproducing the CEB paper’s compression–robustness gradient under a protocol the deterministic path cannot game. Within CEB’s operating range the two geometries are comparable at matched β (at $\beta = 0.1$ under L_∞ : 37.4% for both CEB and OPB, $a = 12$), but the β axis again hides the compression gap — at $\beta = 0.1$ CEB is the *more* compressed of the two (realized $\mathbb{E}[\text{KL}] = 0.68$ vs. 1.23 for OPB, $a = 12$). At compression-matched points the geometry pays: OPB ($a = 12$, $\beta = 1$; $\mathbb{E}[\text{KL}] = 0.55$) attains 30.6% vs. CEB’s 37.4% under L_∞ . The decisive margin is the operating range: OPB keeps accuracy at ≈ 0.98 out to $\beta = 25$ while its targeted success keeps falling, to 11.1%/28.0%/29.5% ($a = 1/6/12$) under L_∞ and 11.0%/33.5%/34.7% under L_2 — 30–50 points below the deterministic baseline and below anything CEB attains with its accuracy intact. At the extreme end the anchor ordering tracks the realized compression: the smallest anchor is also the most collapsed at $\beta = 25$ ($\mathbb{E}[\text{KL}] = 0.013$ vs. 0.28–0.40) and the most robust (11% vs. $\approx 31\%$), the same compression–robustness gradient read along the anchor axis. We do not over-claim: this is one dataset and one target class.

Transfer attacks (gradient-masking check). A white-box attack could fail simply because the target model’s gradients are degenerate, crediting dead models with robustness. We follow the CEB paper’s transfer protocol as the check: generate targeted PGD adversaries on one model (the same 20-step, EOT-10 recipe) and evaluate them on another (10-sample MC predictions), sweeping ε over 0.05–0.3 (L_∞) and 1–6 (L_2). The models are CEB ($\beta = 0.1$, the last point of its operating range), OPB at maximum compression ($\beta = 25$, $a = 1$; realized $\mathbb{E}[\text{KL}] = 0.013$), and the deterministic MLP baseline, whose white-box curve reaching 61% at $\varepsilon_\infty = 0.3$ confirms that the attack itself is not broken.

Figure 3 reads each transfer curve against its target’s own white-box attack, not against the source. CEB→OPB tops out at 6.4% (L_∞) and 6.5% (L_2) — at or below OPB’s own white-box attack (12.0%/12.6%) — so the extreme-compression robustness of OPB survives adversaries that never touch its gradients: it is geometry, not gradient masking. The reverse direction transfers almost nothing ($\leq 1.8\%$), because the collapsed posterior’s perturbations carry little transferable signal — consistent with the reference collapse of CEB being a genuine failure mode rather than a robustness channel.

3.5 MECHANISM VALIDATION: THE GEOMETRY DRIVES REPRESENTATION AND PREDICTION

The main sweep established the predictive adequacy of the constructed geometry; here we verify the mechanism behind it, on the two representative settings of the analysis, MNIST (MLP, classification)




Figure 3: Transfer attacks (gradient-masking check) on MNIST (MLP backbone): targeted PGD adversaries generated on one model and evaluated on another, against ϵ . (a) L_∞ ; (b) L_2 . Solid: white-box self-attacks of CEB ($\beta = 0.1$), OPB ($\beta = 25$, $a = 1$) and the deterministic MLP baseline. Dashed: transfer attacks. CEB→OPB stays at or below OPB’s own white-box attack, so OPB’s extreme-compression robustness is not gradient masking; OPB→CEB transfers almost nothing, as the collapsed posterior’s perturbations carry little transferable signal.

Table 2: Mechanism validation: the constructed geometry pulls the posterior into the orthogonal frame (classification) and onto the isometric axis (regression). Geometric concentration: mean off-diagonal cosine of the posterior class-center Gram matrix (classification) / off-axis residual fraction (regression). Scale following: class-center radius $\|\bar{c}\|$ against the anchor scale a (classification) / axial slope against ρ (regression). Cross-seed stability: std. over runs. Means $\pm$ std. over five runs.

Model (task)	Task	Geometric concentration	Scale following	Cross-seed stability
CEB (MNIST)	Classif.	0.1205 ± 0.0026	5.14	0.0201
OPB (MNIST)	Classif.	0.0203 ± 0.0021	11.00 ($a = 12$)	0.0082
CEB (Housing)	Regress.	0.0233	2.65 ± 0.59	0.593
EPB (Housing)	Regress.	0.0002	9.22 ± 0.15 ($\rho = 12$)	0.155

and California Housing (MLP, regression), trained at the matched compression $\beta = 0.1$ with anchor scale $a = \rho = 12$.

The prior. CEB’s reference is arbitrary and drifts across seeds: the mean off-diagonal cosine of its learned class-prior table is 0.119 ± 0.002 , and the scale of its learned regression axis varies by 2.07 ± 0.88 across runs. The constructed prior has no such freedom: the OPB class priors are exactly orthogonal (off-diagonal cosine 0, pairwise distance $\sqrt{2}a = 16.97$), and the EPB axis is fixed at $\rho = 12.000$ with exact isometry through the origin, both verified at machine precision.

The posterior. Table 2 shows that this geometry transfers to the posterior. On classification the posterior class centers are nearly orthogonal (the mean off-diagonal Gram cosine is $5.9\times$ lower than CEB’s) and sit at radius $\|\bar{c}\| = 11.00$ (92% of a), with $2.5\times$ lower cross-seed drift; the Fisher ratio is comparable (15.3 vs. 15.7, reported as is). On regression the posterior lives essentially *on* the isometric axis: the off-axis residual fraction is $141\times$ lower than CEB’s, and the axial slope tracks ρ with $3.8\times$ lower cross-seed variation. The slope is 9.22 rather than $\rho = 12$ because the encoder’s uncertainty shrinks the posterior along the axis (the slope equals $\rho \cdot R_{\text{axis}}^2$, as $12 \times 0.78 \approx 9.4$), and the axis *direction* itself can rotate rigidly with the representation (mean direction cosine 0.03–0.06 for both models), so the discriminative stability evidence is the scale, not the direction. Figures 4 and 5 show both mechanisms directly.

Prediction reads off the geometry. Both heads of Section 2.2 have no free parameters. At the matched configuration the anchor-energy classifier attains 98.44% test accuracy against CEB’s 98.16% on MNIST, and the tied projection head attains $R^2 = 0.778$ against CEB’s 0.491 on Housing: under this strong compression CEB collapses while the geometric readout retains the signal, consistent with the main-table GPB column (Section 3.2) and with the nearest-anchor analysis of the appendix. The mechanism chain is therefore complete: the constructed prior pulls the posterior




Figure 4: Pairwise-cosine Gram matrix of the posterior class centers on MNIST, averaged over five runs. OPB is nearly diagonal: the posterior centers are orthogonal; CEB shows no such structure.




Figure 5: The axial coordinate $u^\top \mu_q(x)$ against the normalized label on Housing (one representative run per model). The EPB posterior hugs the prior line $\rho \tilde{y}$, shrunk toward $\rho \mathbb{E}[\tilde{y} | x]$ by the encoder’s uncertainty.

into the declared geometry, and a zero-parameter rule that reads exactly that geometry carries the prediction.

4 RELATED WORK

AI USE STATEMENT

(This section is **required** and does not count toward the page limit.)

In this work, we used generative AI tools for drafting and editing portions of the code and of this manuscript. All AI-assisted content was manually reviewed and verified by the authors; experimental results and theoretical claims were produced and checked by our own code and analysis. We take responsibility for the final content of this work, including text, claims or artifacts produced with the aid of generative AI.

REPRODUCIBILITY STATEMENT

The code, hyperparameter grids, and evaluation protocol are described in Section 3.1; the geometric prior is stateless (no EMA or prototypes) and the orthogonalization is exact wherever defined (Assumption 2.1). All datasets are public, with preprocessing steps given in the release. Proofs of all statements are collected in the appendix. Statistical claims use the paired bootstrap confidence intervals and non-inferiority margin described in Appendix A.6.

ACKNOWLEDGMENTS

REFERENCES

- Eslam Abdelaleem, Ibrahima Ndiour, and Abhishek Bhattacharyya. Deep variational multivariate information bottleneck. *Journal of Machine Learning Research (JMLR)*, 2025.
- Alexander A. Alemi, Ian Fischer, Joshua V. Dillon, and Kevin Murphy. Deep variational information bottleneck. In *International Conference on Learning Representations (ICLR)*, 2017.
- Ian Fischer. The conditional entropy bottleneck. *Entropy*, 22(9):999, 2020.
- Artemy Kolchinsky and Brendan D. Tracey. Caveats for information bottleneck in deterministic scenarios. In *International Conference on Learning Representations (ICLR)*, 2019.
- Artemy Kolchinsky, Brendan D. Tracey, and David H. Wolpert. Nonlinear information bottleneck. *Entropy*, 21(12):1181, 2019.
- Naftali Tishby, Fernando C. Pereira, and William Bialek. The information bottleneck method. In *Proc. 37th Allerton Conference on Communication, Control, and Computing*, 2000.

A DEFERRED PROOFS, CONSTRUCTION DETAILS, AND SCOPE

A.1 FAILURE DIAGNOSIS

Proposition A.1 (Dead knob, following Caveat 1 of Kolchinsky & Tracey (2019)). *Assume $Y = f(X)$. For every $\gamma > 0$, the CEB objective $L_\gamma(Z) := I(X; Z|Y) - \gamma I(Y; Z)$ is minimized exactly on the corner family $\{Z : I(Y; Z) = H(Y), I(X; Z|Y) = 0\}$, and the map $\gamma \mapsto (I(X; Z^*|Y), I(Y; Z^*))$ is constant. Consequently no interior point of the IB curve is ever a Lagrangian minimizer of CEB, over the entire parameter range.*

Proof sketch (see Kolchinsky & Tracey (2019) for the full argument). By the chain rule, $L_\gamma(Z) = I(X; Z) - (1 + \gamma)I(Y; Z)$; the data-processing inequality gives $L_\gamma \geq -\gamma H(Y)$, attained exactly at the corner $Z = Y$. In the Caveats paper’s maximizing convention this is their Eq. 7 with $\beta' = 1/(1 + \gamma) \in (0, 1)$, the range in which every maximizer is pinned at the copy corner. $\square$

The corner is a single point of the information plane and an entire equivalence class of representations; the objective is exactly indifferent within the class, and neither the objective nor the certificate records which member the optimizer reaches. In the deterministic scenario this corner *is* the MNI point, so CEB’s declared target and the degenerate attractor coincide. The Caveats paper repairs this by squaring the compression term; our anchored reference supplies the same strict convexity by a different mechanism, without squaring: a quadratic displacement term built into the KL by construction (Proposition A.4 below). We present this as a mechanism connection to the Caveats fix, not as a new solution of the deterministic degeneration.

Proof of Proposition 2.1. For any matched pair ($b = q(z|y)$), the gap identity of Eq. equation 1 gives Eq. equation 2; the matching condition is stated per class and no term of the certificate couples two distinct class priors, so the certificate value carries no information about the pairwise geometry of $\{q(z|y)\}$. Note that $I(X; Z|Y)$ is generally non-zero; it vanishes only under the conditional independence $Z \perp X | Y$, a property of the encoder, not of the prior. $\square$

Proposition A.2 (Retention blindness on T_α). *On the manifold $T_\alpha = B_\alpha \cdot Y$ of Kolchinsky & Tracey (2019) (their Eq. 6), a family that depends on Y alone, $I(X; T_\alpha | Y) = 0$ for every α , and with the optimal backward encoder and classifier, $\text{ReX}(T_\alpha) = 0$ and $\text{CE}(T_\alpha) = H(Y) - I(Y; T_\alpha) =: \delta_\alpha$: the certificate reads 0 from the total collapse ($\alpha = 0$) to the full copy ($\alpha = 1$), while the retained label information $I(Y; T_\alpha)$ sweeps the whole axis. Tightness of the certificate is therefore compatible with retaining all, part, or none of Y : the certificate records nothing about how much label information the code retains.*

Proof. T_α is a function of Y alone, hence $T_\alpha \perp X \mid Y$ and $I(X; T_\alpha | Y) = 0$. By Proposition 2.1 with the true backward encoder the certificate is tight, $\text{ReX} = I(X; T_\alpha | Y) = 0$, and the optimal classifier attains $\langle \log c(y|z) \rangle = -H(Y|T_\alpha)$, i.e. $\text{CE} = H(Y|T_\alpha) = \delta_\alpha$. $\square$

Remark A.1 (The objective on T_α is β -independent and strictly prefers the full copy). On T_α the compression term contributes nothing ($\text{ReX} \equiv 0$), so the objective reads $L = \text{CE} + \beta \text{ReX} = \delta_\alpha$, independent of β and strictly minimized at $\alpha = 1$ (the full copy). No certified residual is produced anywhere on the manifold; the blindness is that the certificate cannot tell the optimizer — or the analyst — where on the retention axis the code sits.

Scope of the deterministic comparison-solution argument. In Theorem 2.1 the comparison solution $q(z|x) = p_\theta(z|y)$ requires a zero-KL feasible solution. For tasks in which none exists (in particular for continuous regression, where the encoder cannot recover y from x and no such Bayes comparison classifier is defined), part (a) does not apply: there the separation transfer of Proposition 2.4 remains a conditional statement on the anchor mismatch, and no framework-level alignment bound is claimed. Part (b) is a separate deterministic-regression argument with its own comparison solution.

A.2 PROOFS OF THE GUARANTEES

Proof of Proposition 2.3. For fixed θ , apply the gap identity of Eq. equation 1 to the conditional distribution $p_\theta(z|y)$ and the posterior q_ϕ . $\square$

Proof of Proposition 2.4. For (a), Jensen gives $\|m_k - m_\theta(k)\| \leq \sqrt{2\tau_{\max}^2 D_k}$, and the triangle inequality on $m_i - m_j = (m_i - m_\theta(i)) + (m_\theta(i) - m_\theta(j)) + (m_\theta(j) - m_j)$ gives Eq. equation 16. For (b), the same two inequalities hold pointwise for P_Y -a.e. y : Jensen on the regular conditional expectation gives $\|m(y) - m_\theta(y)\| \leq \sqrt{2\tau_{\max}^2 D(y)}$ (the norm is convex), and the triangle inequality gives Eq. equation 18; intersecting the two measure-one sets of valid y 's yields the almost-everywhere pair statement. $\square$

Remark A.2 (Finite-sample estimation of the continuous case). On finite samples the regular conditional expectations of part (b) are estimated by binning the label axis into intervals B_i with representative labels $\bar{y}_i$ and centers $m_i = \mathbb{E}[\mu_\phi(X) \mid Y \in B_i]$; the bin diameter adds a quantization error bounded by $\rho_{\max} \text{diam}(B_i) + \rho_{\max} \text{diam}(B_j)$ (by the upper bound of Eq. equation 5), which vanishes as the bins shrink. The theorem itself is stated pointwise and needs no binning.

Proof of Theorem 2.1. (a) Since $\Sigma_\theta(y) \preceq \tau_{\max}^2 I_d$, each sample's KL contains a mean-mismatch term with variance at most $\tau_{\max}^2$, so $\text{KL}(q_\phi(z|x) \parallel p_\theta(z|y)) \geq \|\mu_\phi(x) - m_\theta(y)\|^2 / (2\tau_{\max}^2)$, and taking expectations gives $\mathbb{E}[\text{KL}] \geq \mathbb{E}[D(Y)]$. The comparison solution $q(z|x) = p_\theta(z|y)$ is feasible by the assumptions ($Y = f(X)$, class-constant maps, diagonal representable covariances), has zero KL, and its prediction is the Bayes posterior of Eq. equation 19, whose cross-entropy is exactly the conditional entropy $C_G = H(Y \mid Z_{\text{align}}) \leq \ln K$; for the orthogonal instance the anchor-energy classifier of the main model realizes it (shared covariance, determinant factors cancel). Hence $L^* \leq C_G$ and $\beta \mathbb{E}[\text{KL}] \leq L \leq L^* + \varepsilon \leq C_G + \varepsilon \leq \ln K + \varepsilon$, whence Eq. equation 20.

(b) As in (a), $\mathbb{E}[\text{KL}] \geq \mathbb{E}[D(Y)]$ by the covariance bound. The comparison solution is feasible by the assumptions ($Y = f(X)$, aligned maps representable): its posterior is $q(z|x) = p_\theta(z|y)$ (zero KL) and its prediction is the tied head, whose risk is $C_{\text{align}} = \tau^2 / \rho^2$ (the prediction error is exactly the Gaussian bottleneck noise $(\tau/\rho)\eta$). Hence $L^* \leq C_{\text{align}}$ and $\beta \mathbb{E}[\text{KL}] \leq L \leq L^* + \varepsilon \leq C_{\text{align}} + \varepsilon$, whence Eq. equation 21. $\square$

Proof of Corollary 2.1. (a) $D_k \leq \mathbb{E}[D(Y)]/\pi_k \leq (C_G + \varepsilon)/(\beta\pi_k)$; combine with Proposition 2.4, whose mismatch terms are $\sqrt{2\tau_{\max}^2 D_k}$. (b) Markov’s inequality on the non-negative variable $D(Y)$ with mean at most $(C_{\text{align}} + \varepsilon)/\beta$ gives $P_Y(D(Y) > t) \leq (C_{\text{align}} + \varepsilon)/(\beta t)$; on S_t the pointwise transfer of Proposition 2.4(b) applies with $D(y), D(y') \leq t$. $\square$

Proof of Proposition 2.5. By Jensen’s inequality the left side is at least $\frac{1}{2\tau^2} \sum_k \pi_k \|c - a q_{\theta,k}\|^2 = \frac{1}{2\tau^2} (\|c\|^2 - 2a c^\top \sum_k \pi_k q_{\theta,k} + a^2)$. Minimizing over c gives $c = a \sum_k \pi_k q_{\theta,k}$; by orthonormality $\|\sum_k \pi_k q_{\theta,k}\|^2 = \sum_k \pi_k^2$, yielding Eq. equation 24. $\square$

Remark A.3 (What the framework constrains (and what it does not)). The constraints of the geometric prior bind the reference, not the deployed representation: for any orthonormal frame $\{q'_k\}$ and the class-constant posterior $q(z|Y = k) = \mathcal{N}(a q'_k, \tau^2 I_d)$, the trainable frame can rotate to match $(q_{\theta,k} \rightarrow q'_k)$, yielding $\text{KL} = 0$, $I(X; Z|Y) = 0$, and $I(Y; Z) > 0$: a full-retention label code with a zero certificate. The framework therefore constrains the geometric *form* of the label code (precisely what CEB left unspecified), and does not claim that the certificate meters retained label information. (With the variance fixed, the remaining freedom at certificate zero is the frame rotation; a learnable per-class variance would add a further circumvention channel and is deliberately not used.)

A.3 CONSTRUCTION DETAILS

Polar permutation equivariance. For any label permutation matrix P ,

$$Q_\theta(U_\theta P) = U_\theta P (P^\top U_\theta^\top U_\theta P)^{-1/2} = U_\theta (U_\theta^\top U_\theta)^{-1/2} P = Q_\theta(U_\theta) P, \quad (26)$$

since $PP^\top = I_K$ and the inverse square root conjugates under P . Permuting the label numbering therefore permutes the anchors accordingly, and no class index plays a distinguished role.

Proposition A.3 (Orthogonal class-prior geometry by construction). *Under Eqs. equation 8–equation 9, the class means $a_k = a q_{\theta,k}$ satisfy Eq. equation 10; equivalently, the class-mean Gram matrix is $a^2 I_K$: under Assumption 2.1, the prior supplies a class-separated spherical code at every training step, whatever the optimizer does.*

Proof. Eq. equation 10 follows from $Q_\theta^\top Q_\theta = I_K$ after scaling by a ; in particular $\|a_i - a_j\|^2 = \|a_i\|^2 + \|a_j\|^2 - 2a_i^\top a_j = 2a^2$. $\square$

Initialization, monitoring, and re-initialization. The posterior logvar head is zero-initialized ($\sigma^2 = \tau^2 = 1$ at initialization, so training starts at $\text{KL} \approx \frac{1}{2\tau^2} a^2$); the prior encoder is identity-initialized, so the initial raw matrix is $U_\theta = [e_1, \dots, e_K]$, full column rank, with a stable polar-factor backward (the verified condition of Assumption 2.1). The full-rank condition is *verified, not imposed*: $\sigma_{\min}(U_\theta)$ is monitored along training and stays bounded away from zero in the reported runs, and no additional regularizer is added to the objective, which remains Eq. equation 11 exactly. Near rank deficiency the polar factor’s gradient is delicate; should the monitor flag a near-deficient step, the recommended handling is a re-initialization of the prior encoder (the reported runs never triggered it).

EPB axis choices. Three variants of the direction: (A) *fixed axis*, $u = [1, 0, \dots, 0]$, with no trainable prior direction: the theoretically cleanest ablation, at the price of prespecifying a latent coordinate; (B) *trainable normalized axis*, $u = W/\|W\|$ with W trained and normalized at every step, our main variant, which keeps the scale ρ exact while allowing the model to choose the axis; (C) *unconstrained linear axis*, $\mu_p(y) = W\tilde{y}$: the simplest ablation, still isometric by linearity but with a free scale that can shrink and weaken the geometry. We recommend (B), use (A) to test whether the trainable rotation carries the benefit, and keep (C) as a diagnostic. The main variant uses no bias: an affine offset cancels in pairwise distances but is a free extra degree of matching, and $\tilde{y}$ is already centered.

EPB isometry notes. The isometry is exact and global: any nonzero linear map of a scalar label is isometric up to its norm, and the normalization fixes that norm to ρ ; unlike fixed random-feature anchors, whose pairwise distances grow nonlinearly and saturate, the isometric axis preserves label distances at every scale (with the price that anchor norms grow with $|\tilde{y}|$; a line and a fixed-radius

sphere intersect in at most two points, so isometry and equal norms cannot both hold). All prior means lie on one line; for $y_i < y_j < y_k$ the distances add, $\|\mu_p(y_i) - \mu_p(y_k)\| = \|\mu_p(y_i) - \mu_p(y_j)\| + \|\mu_p(y_j) - \mu_p(y_k)\|$: order and magnitude of label differences are encoded exactly.

Multidimensional regression labels. For $y \in \mathbb{R}^m$ ($m \leq d$), the same construction extends verbatim: take $W \in \mathbb{R}^{d \times m}$, its thin polar factor $Q = W(W^\top W)^{-1/2}$, and the prior $\mu_p(y) = \rho Q \tilde{y}$, which is isometric in the label metric; the scalar case is $m = 1$ with $Q = u$. The label metric must be declared before standardization, otherwise “isometric” has no unique meaning.

A.4 ADDITIONAL RESULTS

Lemma A.1 (OPB KL decomposition). *For a diagonal posterior covariance and the prior of Eq. equation 9,*

$$\text{KL}(q_\phi(z|x) \| p_\theta^{\text{OPB}}(z|y)) = \frac{1}{2\tau^2} \|\mu_\phi(x) - a q_{\theta,y}\|^2 + \frac{1}{2} \sum_{r=1}^d \left[\frac{\sigma_{\phi,r}^2(x)}{\tau^2} - 1 - \log \frac{\sigma_{\phi,r}^2(x)}{\tau^2} \right], \quad (27)$$

a class-anchor mismatch term plus a variance mismatch term that is non-negative and vanishes exactly at $\sigma_{\phi,r}^2(x) = \tau^2$ for all r . The mismatch term is quadratic in the displacement to a reference whose scale and covariance cannot move: this is the mechanism behind the strict convexity of Proposition A.4 below.

Proof. Apply the diagonal-Gaussian KL formula with posterior mean $\mu_\phi(x)$, prior mean $a q_{\theta,y}$, posterior variances $\sigma_{\phi,r}^2(x)$, and the fixed prior variance τ^2 . $\square$

Corollary A.1 (Label-information lower bound in the aligned Gaussian model). *Assume uniform classes and exact alignment $q(z|Y = k) = \mathcal{N}(a q_{\theta,k}, \tau^2 I_d)$. The pairwise error probability of the nearest-anchor classifier is $P_{\text{pair}} = \Phi(-\frac{a}{\sqrt{2}\tau})$, with $P_e \leq \min\{1, (K-1)P_{\text{pair}}\}$, and whenever the displayed bound $\bar{P}_e \leq 1 - 1/K$, Fano’s inequality gives $I(Y; Z) \geq \log K - h_2(\bar{P}_e) - \bar{P}_e \log(K-1)$. This links the anchor scale-to-noise ratio a/τ to the retained label information, conditionally on the aligned Gaussian model.*

Proof. Two anchors have distance $\sqrt{2}a$ by Proposition A.3; the equal-covariance binary decision boundary is halfway between them, giving $\Phi(-\sqrt{2}a/(2\tau))$. The union bound and Fano’s inequality on $[0, 1 - 1/K]$ conclude. $\square$

Proposition A.4 (Joint strict convexity and unique optimum on a restricted aligned surrogate). *Consider the class-symmetric aligned family: every input of class y is mapped to the posterior $q_y = \mathcal{N}(s q_{\theta,y}, \sigma^2 I)$ ($s \geq 0$, $\sigma^2 > 0$; anchor scale $a > 0$), and the classifier is the anchor-energy classifier of Proposition 2.2 (weights $w_j = c q_{\theta,j}$ with scale $c = a/\tau^2$ fixed by construction). With the stochastic training code $z \sim q_y$ (as in the actual objective, Eq. equation 11), the training cross-entropy is*

$$g(s, \sigma^2) := \mathbb{E}_\eta \left[\ln \left(1 + \sum_{j \neq y} e^{-cs + c\sigma(\eta_j - \eta_y)} \right) \right], \quad \eta \sim \mathcal{N}(0, I_K), \quad (28)$$

and the training Lagrangian is

$$L_\beta(s, \sigma^2) = g(s, \sigma^2) + \beta \left[\frac{1}{2\tau^2} (s - a)^2 + \frac{d}{2} (\sigma^2/\tau^2 - 1 - \ln(\sigma^2/\tau^2)) \right]. \quad (29)$$

Then for every $\beta > 0$, L_β is strictly convex jointly in (s, σ) and has a unique joint minimizer $(s^(\beta), \sigma^*(\beta))$, interior in s since $a > 0$: the stochastic training objective has a well-defined optimum, in contrast to the deterministic IB Lagrangian, whose optimum is pinned at a corner for every β (Proposition A.1). This is a restricted surrogate result: it analyzes two scalar variables on the class-symmetric aligned family, with the frame, the anchor scale, and the classifier scale held fixed; the full OPB objective, with its nonlinear posterior, trainable frame, and non-convex degrees of freedom, lies outside this analysis, and the result neither establishes β -tunability of the full objective nor makes any claim on the information plane. The deterministic evaluation code ($z = \mu$) is the $\sigma \rightarrow 0$ limit, and $g \geq \ln(1 + (K-1)e^{-cs})$ by Jensen (the log-sum-exp is convex), so the*

statement strictly generalizes the deterministic surrogate analysis. We do not claim here that the map $\beta \mapsto s^*(\beta)$ is monotone or that the sweep is one-to-one: the minimizer σ^* responds to β through the cross-derivative $\partial^2 g / \partial s \partial \sigma$, and establishing the sign of $ds^* / d\beta$ requires the full two-dimensional stationarity analysis, which we leave open; the β -sweep is reported empirically in the experiments. The fixed-scale condition holds for the main model by construction (Proposition 2.2), so the statement describes the actual training objective up to the aligned-family idealization. We do not analyze the free-classifier variant on this surrogate: under the stochastic training code, a diverging classifier scale does not produce a degenerate solution: at $KL \rightarrow 0$ the posterior variance is pinned at $\sigma^2 = \tau^2 > 0$, so the class-conditional Gaussians overlap and the training cross-entropy of a free classifier diverges with its scale rather than vanishing (for $K = 2$ it contains $\mathbb{E}[\ln(1 + e^{-c(a+\tau\xi)})]$, whose integrand grows linearly in c on the event $\xi < -a/\tau$, which has positive probability). An analysis of the free classifier would require an explicit classifier-norm constraint or regularizer and is left to future work.

Proof. For each fixed η , the integrand of Eq. equation 28 is a log-sum-exp of functions affine in (s, σ) (the constant 1 has slope zero and each exponential has slope $(-c, c(\eta_j - \eta_y))$), hence jointly convex in (s, σ) , and the expectation g is jointly convex. No strictness is claimed for g : for $K = 2$ the integrand depends on (s, σ) only through the single linear combination $-cs + c\sigma(\eta_2 - \eta_1)$, so its Hessian has rank one. Write the KL term as $R(s, \sigma) := \frac{1}{2\tau^2}(s - a)^2 + \frac{d}{2}(\sigma^2/\tau^2 - 1 - \ln(\sigma^2/\tau^2))$; its Hessian in (s, σ) is $\nabla^2 R = \text{diag}(\tau^{-2}, d(\tau^{-2} + \sigma^{-2}))$, strictly positive definite for every $\sigma > 0$, so βR is strictly convex in (s, σ) . Hence $L_\beta = g + \beta R$ is the sum of a jointly convex and a strictly convex function, and is strictly convex in (s, σ) . Since $R \rightarrow +\infty$ as $s \rightarrow \infty$, $\sigma \rightarrow 0^+$, or $\sigma \rightarrow \infty$, the sublevel sets of L_β are compact, and the unique joint minimizer exists (by strict convexity it is unique). For $a > 0$ the minimizer is interior in s : at $s = 0$, $\partial L_\beta / \partial s = -c\mathbb{E}[p_w] - \beta a / \tau^2 < 0$, where p_w is the error probability of the sampled code, so $s^* > 0$. $\square$

Remark A.4 (The sweep is a KL–CE surrogate curve; scope and limitations). Proposition A.4 proves joint strict convexity and uniqueness of the optimum for the stochastic training objective in the KL–CE plane on the aligned family at fixed classifier scale: a restricted surrogate result, which does not establish β -tunability of the full OPB objective, and in any case not a statement on the information plane of the Caveats paper; the results of the main text establish a structured conditional prior, an exact KL geometry, and conditional separation guarantees, and they do *not* establish a mutual information upper bound tighter than CEB’s: Eq. equation 14 is the usual CEB gap identity for a particular prior family. The orthogonal-frame construction fixes all pairwise anchor distances, so the framework should not claim to recover arbitrary semantic class similarities. All geometric statements are conditional on Assumption 2.1: orthonormality holds exactly by the polar construction wherever the factorization exists, and the condition is verified, not imposed, along the reported runs (Section 2.2.1); the construction is exactly equivariant under relabeling of the classes by Eq. equation 26, so no class index plays a distinguished role. If the frame were computed from the current input batch rather than from the all-class prior matrix, the prior would become batch-dependent and Proposition 2.3 would no longer follow directly. Finally, the anchor-energy classifier of Proposition 2.2 fixes the classifier scale to a/τ^2 by construction, removing the independent classifier-margin escape route in the aligned surrogate of Proposition A.4; whether a free classifier (with an explicit norm constraint or regularizer) admits its own degenerate route is left to future work, and the free linear classifier is kept in the experiments as a variant for comparison. Whether the joint optimum of Proposition A.4 moves monotonically in β is likewise open: it requires the full two-dimensional stationarity analysis, in which the optimal variance responds to β through the cross-derivative $\partial^2 g / \partial s \partial \sigma$ of the stochastic cross-entropy. The trainable frame and axis constitute a rotation channel that this paper does not address: the certificate’s zero set contains the full-retention orthogonal label codes at any frame rotation (the reference can rotate to meet them), and the framework pins the *geometry of the code*; it does not claim that the certificate reads off the retained label information. Quantifying how much of the anchor gap is absorbed by frame rotation rather than by posterior motion is left to future work (a frozen-frame ablation isolates it).

A.5 MAIN RESULTS WITH STANDARD DEVIATIONS

Tables 3 and 4 report the same configurations as the main table with standard deviations over runs, split by column groups for width; bold and underline follow the same mean-based rule as Table 1.

Table 3: Test accuracy (%) and test R^2 with standard deviations over runs: the plain backbone and the variational baselines; same configurations as Table 1.

Dataset	Backbone	Base	VIB	SVIB	NIB
MNIST	MLP	98.33±0.12	98.28±0.11	98.15±0.13	98.27±0.07
MNIST	CNN	99.13±0.05	99.28±0.09	99.23±0.03	99.22±0.06
ImageNet-100	MLP	83.12±0.14	82.84±0.44	83.02±0.37	82.35±0.27
ImageNet-100	CNN	83.40±0.13	82.94±0.37	83.11±0.86	82.91±0.35
Cora	GCN	87.68±0.46	86.83±0.28	86.46±1.87	86.35±0.68
IMDb	RNN	84.85±0.68	85.50±0.91	85.99±0.30	86.14±0.92
AG News	RNN	92.94±0.15	92.72±0.17	92.86±0.21	92.91±0.22
Cal. Housing	MLP	0.7618±0.0238	0.7901±0.0049	0.7897±0.0058	0.7862±0.0035
STS-B	RNN	0.2899±0.0113	0.2719±0.0164	0.2697±0.0202	0.2417±0.0221
ZINC	GCN	0.6348±0.0044	0.6339±0.0033	0.6325±0.0026	0.6178±0.0099
AgeDB	MLP	0.2180±0.0409	0.3173±0.0101	0.3385±0.0286	0.2710±0.0235
AgeDB	CNN	0.3914±0.0194	0.3342±0.1894	0.3323±0.1870	0.3337±0.1889

Table 4: Test accuracy (%) and test R^2 with standard deviations over runs: CEB, DVCCA, and our method; same configurations as Table 1. GPB and GPB-L are separated from the references by the vertical rule.

Dataset	Backbone	CEB	DVCCA	GPB	GPB-L
MNIST	MLP	98.43±0.08	98.24±0.17	98.54±0.09	98.43±0.07
MNIST	CNN	99.30±0.05	99.29±0.10	99.32±0.05	99.31±0.05
ImageNet-100	MLP	82.77±0.22	82.99±0.43	84.40±0.21	84.01±0.23
ImageNet-100	CNN	82.97±0.67	83.21±0.70	86.08±0.21	84.88±0.80
Cora	GCN	86.27±0.38	86.90±0.94	88.12±0.55	87.86±0.87
IMDb	RNN	85.55±0.73	85.69±0.23	86.84±0.57	86.99±0.35
AG News	RNN	92.83±0.11	92.78±0.15	92.87±0.12	92.92±0.14
Cal. Housing	MLP	0.7894±0.0018	0.7878±0.0048	0.7895±0.0051	0.7914±0.0034
STS-B	RNN	0.2668±0.0127	0.2713±0.0218	0.2762±0.0152	0.2755±0.0134
ZINC	GCN	0.6257±0.0044	0.6294±0.0026	0.6344±0.0031	0.6323±0.0016
AgeDB	MLP	0.3317±0.0251	0.3180±0.0228	0.3536±0.0076	0.3320±0.0080
AgeDB	CNN	0.3882±0.0275	0.3205±0.1811	0.4700±0.0072	0.4388±0.0220

A.6 PAIRED-DIFFERENCE CONFIDENCE INTERVALS

Tables 5–7 report, for each setting, the paired difference of GPB over every baseline (GPB minus baseline) with a bootstrap 95% confidence interval: the differences are paired across the shared seeds, the interval is the percentile interval of $B = 10,000$ resamples of the per-run differences (fixed seed, reproducible), and the non-inferiority margin is $\delta = 0.2$ percentage points of accuracy or $\delta = 0.005$ in R^2 . Bold marks a lower bound above zero (significant advantage) and † a lower bound within the margin (tied); the best configuration per method is the same as in the main table.

Table 5: Paired-difference 95% confidence intervals of GPB over the plain backbone and VIB (GPB minus baseline, per setting; classification in percentage points, regression in R^2). Bold: lower CI bound above zero; $\dagger$: lower bound within the non-inferiority margin δ (tied).

Dataset	Backbone	Base	VIB
MNIST	MLP	+0.2 [+0.2, +0.3]	+0.3 [+0.2, +0.3]
MNIST	CNN	+0.2 [+0.2, +0.2]	+0.0 [-0.0, +0.1] $\dagger$
ImageNet-100	MLP	+1.3 [+1.2, +1.4]	+1.6 [+1.4, +1.9]
ImageNet-100	CNN	+2.7 [+2.5, +2.9]	+3.1 [+3.0, +3.4]
Cora	GCN	+0.4 [-0.1, +1.1] $\dagger$	+1.3 [+0.7, +1.7]
IMDb	RNN	+2.0 [+1.2, +2.7]	+1.3 [+0.6, +2.3]
AG News	RNN	-0.1 [-0.2, +0.0] $\dagger$	+0.2 [+0.1, +0.2]
Cal. Housing	MLP	+0.028 [+0.008, +0.048]	-0.001 [-0.005, +0.003] $\dagger$
STS-B	RNN	-0.014 [-0.022, -0.006]	+0.004 [-0.017, +0.022]
ZINC	GCN	-0.000 [-0.003, +0.002] $\dagger$	+0.001 [-0.004, +0.004] $\dagger$
AgeDB	MLP	+0.136 [+0.104, +0.167]	+0.036 [+0.026, +0.046]
AgeDB	CNN	+0.079 [+0.066, +0.091]	+0.136 [+0.043, +0.302]
		8*/3 $\dagger$	8*/3 $\dagger$

Table 6: Paired-difference 95% confidence intervals of GPB over SVIB and NIB (GPB minus baseline, per setting; classification in percentage points, regression in R^2). Bold/ $\dagger$ as in the baselines table.

Dataset	Backbone	SVIB	NIB
MNIST	MLP	+0.4 [+0.3, +0.5]	+0.3 [+0.2, +0.3]
MNIST	CNN	+0.1 [+0.1, +0.1]	+0.1 [+0.1, +0.2]
ImageNet-100	MLP	+1.4 [+1.1, +1.7]	+2.1 [+1.8, +2.2]
ImageNet-100	CNN	+3.0 [+2.4, +3.9]	+3.2 [+2.9, +3.5]
Cora	GCN	+1.7 [+0.2, +3.4]	+1.8 [+1.1, +2.4]
IMDb	RNN	+0.8 [+0.2, +1.4]	+0.7 [+0.2, +1.2]
AG News	RNN	+0.0 [-0.1, +0.1] $\dagger$	-0.0 [-0.2, +0.1] $\dagger$
Cal. Housing	MLP	-0.000 [-0.006, +0.004]	+0.003 [-0.001, +0.009] $\dagger$
STS-B	RNN	+0.006 [-0.021, +0.028]	+0.035 [+0.011, +0.058]
ZINC	GCN	+0.002 [-0.001, +0.004] $\dagger$	+0.017 [+0.010, +0.026]
AgeDB	MLP	+0.015 [-0.010, +0.041]	+0.083 [+0.062, +0.100]
AgeDB	CNN	+0.138 [+0.045, +0.303]	+0.136 [+0.040, +0.302]
		7*/2 $\dagger$	10*/2 $\dagger$

Table 7: Paired-difference 95% confidence intervals of GPB over CEB, DVCCA, and the free-head ablation GPB-L (GPB minus baseline, per setting; classification in percentage points, regression in R^2). Bold/ $\dagger$ as in the baselines table.

Dataset	Backbone	CEB	DVCCA	GPB-L
MNIST	MLP	+0.1 [+0.0, +0.2]	+0.3 [+0.2, +0.4]	+0.1 [+0.0, +0.2]
MNIST	CNN	+0.0 [-0.0, +0.1] $\dagger$	+0.0 [-0.1, +0.1] $\dagger$	+0.0 [-0.0, +0.1] $\dagger$
ImageNet-100	MLP	+1.6 [+1.5, +1.7]	+1.4 [+1.1, +1.7]	+0.4 [+0.2, +0.5]
ImageNet-100	CNN	+3.1 [+2.4, +3.7]	+2.9 [+2.4, +3.5]	+1.2 [+0.5, +1.9]
Cora	GCN	+1.8 [+1.5, +2.3]	+1.2 [+0.4, +1.9]	+0.3 [-0.5, +1.3]
IMDb	RNN	+1.3 [+0.3, +2.1]	+1.1 [+0.8, +1.5]	-0.2 [-0.6, +0.5]
AG News	RNN	+0.0 [-0.1, +0.2] $\dagger$	+0.1 [-0.0, +0.2] $\dagger$	-0.0 [-0.2, +0.1] $\dagger$
Cal. Housing	MLP	+0.000 [-0.003, +0.004] $\dagger$	+0.002 [-0.005, +0.009]	-0.002 [-0.005, +0.002]
STS-B	RNN	+0.009 [-0.009, +0.025]	+0.005 [-0.020, +0.026]	+0.001 [-0.019, +0.017]
ZINC	GCN	+0.009 [+0.005, +0.011]	+0.005 [+0.002, +0.007]	+0.002 [-0.001, +0.005] $\dagger$
AgeDB	MLP	+0.022 [+0.006, +0.038]	+0.036 [+0.012, +0.057]	+0.022 [+0.019, +0.025]
AgeDB	CNN	+0.082 [+0.064, +0.099]	+0.149 [+0.065, +0.310]	+0.031 [+0.018, +0.046]
		8*/3 $\dagger$	8*/2 $\dagger$	5*/3 $\dagger$